TAEHYUN KIM

Dept. of ECE, Seoul National University, Seoul, South Korea
⋆ E-mail: taehyunzzz@snu.ac.kr; kimth@capp.snu.ac.kr; ⋆ Homepage: https://taehyunzzz.github.io/

Research Interests

- **Efficient AI Algorithms**: Parameter-Efficient Finetuning (PEFT), Speculative Decoding, Mixture-of-Experts (MoE)
- **Hardware Architecture**: Processing-in-Memory (PIM), Near-Data Processing (NDP), CXL Memory
- **High-Performance System Design**: Accelerator Parallelism, Algorithm-System Co-design, Memory Optimizations

Work Experience

Computer Architecture and Parallel Processing Lab @ SNU — Mar. 2019 – Present
Ph.D. Researcher — *Seoul, South Korea*

Computer Architecture and Parallel Processing Lab @ SNU — Aug. 2019 – Dec. 2019
Undergraduate Researcher — *Seoul, South Korea*

LG Electronics — Jan. 2018 – Feb. 2018
Undergraduate Intern — *Seoul, South Korea*

Education

Seoul National University (SNU) — Mar. 2019 – Present
Doctor of Philosophy in Electrical and Computer Engineering — *Seoul, South Korea*
- Advisor : Prof. Hyuk-Jae Lee
- GPA: 3.94 / 4.3

Seoul National University (SNU) — Mar. 2013 – Feb. 2019
Bachelor of Science in Electrical and Computer Engineering — *Seoul, South Korea*
- GPA: 3.61 / 4.3

Publications (Featured)

SpecMoE: Memory-Efficient Acceleration of Large Mixture Models with Self-Speculative Decoding
Taehyun Kim, Hyuk-Jae Lee, Jaewoong Sim
Under Review.

MoNDE: Mixture of Near-Data Experts for Large-Scale Sparse Models
Taehyun Kim, Kwanseok Choi, Youngmock Cho, Jaehoon Cho, Hyuk-Jae Lee, Jaewoong Sim
61th ACM/IEEE Design Automation Conference (DAC), Jun. 2024.

An FPGA-based Evaluation Platform for Testing Memory Prototype Chips
Youngmock Cho, Taehyun Kim, and Hyuk-Jae Lee
International Conference on Electronics, Information, and Communication (ICEIC), Jan. 2024.

Analyzing the Scaling Charateristics of Transformer Feed-forward Networks for the Trillion-Parameter Era and Beyond
Taehyun Kim and Hyuk-Jae Lee
International Conference on Electronics, Information, and Communication (ICEIC), Jan. 2024.

Exploring the Inefficiencies of a Large-scale Deep Neural Network Training Framework
Taehyun Kim, Youngmock Cho, and Hyukjae Lee
International Conference on Electronics, Information, and Communication (ICEIC), Feb. 2023.

Virtual Keyboards with Real-time and Robust Deep Learning-based Gesture Recognition
Tae-Ho Lee, Sunwoong Kim, Taehyun Kim, Jin-Sung Kim, Hyuk-Jae Lee
IEEE Transactions on Human-Machine Systems (THMS), Volume 52 Issue 4, Apr. 2022.

Smart Refrigerator Inventory Management using Convolutional Neural Networks
Tae-Ho Lee, Shin-Woo Kang, Taehyun Kim, Jin-Sung Kim, Hyuk-Jae Lee
3rd IEEE International Conference on Artificial Intelligence Circuits and Systems (AICAS), Jun. 2021.

GradPIM: A Practical Processing-in-DRAM Architecture for Gradient Descent
Heesu Kim, Hanmin Park, Taehyun Kim, Kwanheum Cho, Eojin Lee, Soojung Ryu, Hyuk-Jae Lee, Kiyoung Choi, Jinho Lee
27th IEEE International Symposium on High-Performance Computer Architecture (HPCA), Jan. 2021.

Patents

METHOD FOR ROUTING TOKEN AND APPARATUS THEREFOR
Inventor: Taehyun Kim, Jaewoong Sim, Jaehoon Cho, Hyuk-Jae Lee
Submitted: Dec. 10, 2024. (*Under Review*)

NEAR-DATA PROCESSING SYSTEM FOR LARGE-SCALE MIXTURE-OF-EXPERTS AI MODEL
Inventor: Taehyun Kim, Hyuk-Jae Lee, Jaewoong Sim, Jaehun Cho, Kanseok Choi, Youngmock Cho
Submitted: Dec. 29, 2023. (*Under Review*)

Research Projects

Intelligent In-Memory Error-Correction Device for Highly-Reliabile Memory

- Mar. 2021 - Dec. 2024
- Served as the lead manager of the project.
- Developed and verified a hardware architecture of next-generation error-correcting code (ECC) memory.
- Hands-on HDL experience including Verilog design, hardware synthesis and FPGA emulation.
- Conducted research on near-data-processing (NDP) in emerging CXL memory that could work aside ECCs.
- 1 major publication, 2 patents under review.

High-Efficiency Deep Learning based on Memory-Aware Optimizations

- Mar. 2020 - Oct. 2020
- Participated in exploring optimizations for mixed-precision DNN training in an NPU accelerator setting.
- Analyzed the effect of on-chip buffers on energy efficiency and processing speed.
- Designed a PIM accelerator that exploits bank-level parallelism in modern DRAMs to accelerate the memory-intensive parameter update operation in mixed-precision DNN training settings.
- 1 major publication.

Smart Refrigerator Stock Management System

- Jun. 2018 - Dec. 2019
- Developed an automatic stock management system based on YOLOv3.
- Tracks the movement of stored items inside the refrigerator for stock management.
- Exhibited at CES 2020.

Scholarships & Funding

- Samsung, Device Solutions System LSI Division, scholarship contract

Skills & Tools

- Languages (proficiency): Korean (native), English (high), Japanese (high)
- Programming Languages: Python, C++, CUDA, Verilog
- DL Frameworks: PyTorch, HuggingFace, Deepspeed, TensorRT-LLM (FasterTransformer), vLLM
- Simulators: DRAMsim3, Ramulator
- Other tools: SystemC, Synopsys Design Compiler, ModelSim, NVIDIA NSight Systems, Xilinx FPGA, vim

Services

External Reviewer
42nd IEEE International Conference on Computer Design (ICCD) *2024*

Research Assistant
Graduation Project for Undergraduate Students *Spring 2019, Spring & Fall 2024*

Teaching Assistant
400.018 Creative Engineering Design *Fall 2019-2022*

Republic of Korea Army
Administrative staff @ AFMC *2015-2016*